



Hypervisor-agnostic Guest Overlays for VMs

Jörg Thalheim, **Peter Okelmann**, Harshavardhan Unnibhavi,
Redha Gouicem, Pramod Bhatotia



THE UNIVERSITY
of EDINBURGH


Virtual Machines (VMs)

VMs:

- Consolidation
- Cost-effectiveness

Optimized, lightweight VMs:

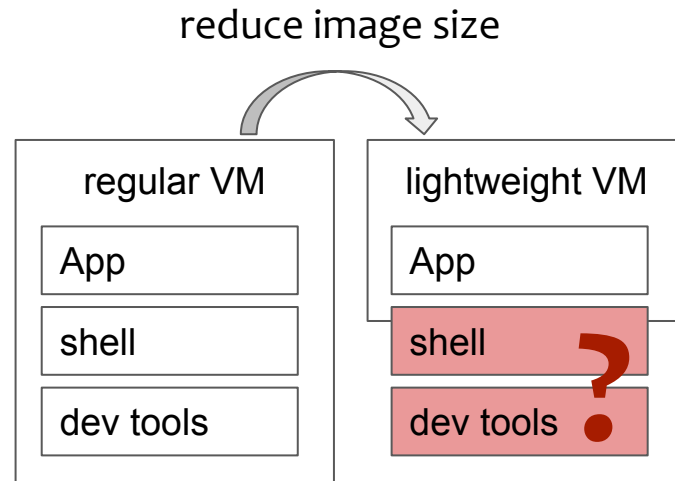
- **Small** memory footprint
- **Fast** bootup times
- Improve **dependability**: trust, reliability



Tradeoff: Lightweight VMs

Limited observability:

- No monitoring and inspection tools
- Disruptive: re-deployment for every change



Debugging, monitoring and repairing is time-consuming

Common solution: VM agents

Agent tasks:

- Provisioning
- Monitoring, Inspection
- Maintenance, Recovery

Multitude of implementations:

Amazon SSM, Google OS Config,
Google Guest Agent, Microsoft OMI,
QEMU Guest Agent, SSH, ...



Overheads for the customer:



Devel & testing:

Provider, Hypervisor and OS distro specific



Infrastructure maintenance:

Management network, key management




Complicated to use:

1600 pages of user manual

VM agents are an unsatisfactory solution


On monolithic servers, providers want to:

- Reduce overheads for customers
- Offer services to customers
 - Out-of-band management (~IPMI)
 - Update notifications
 - Security inspection

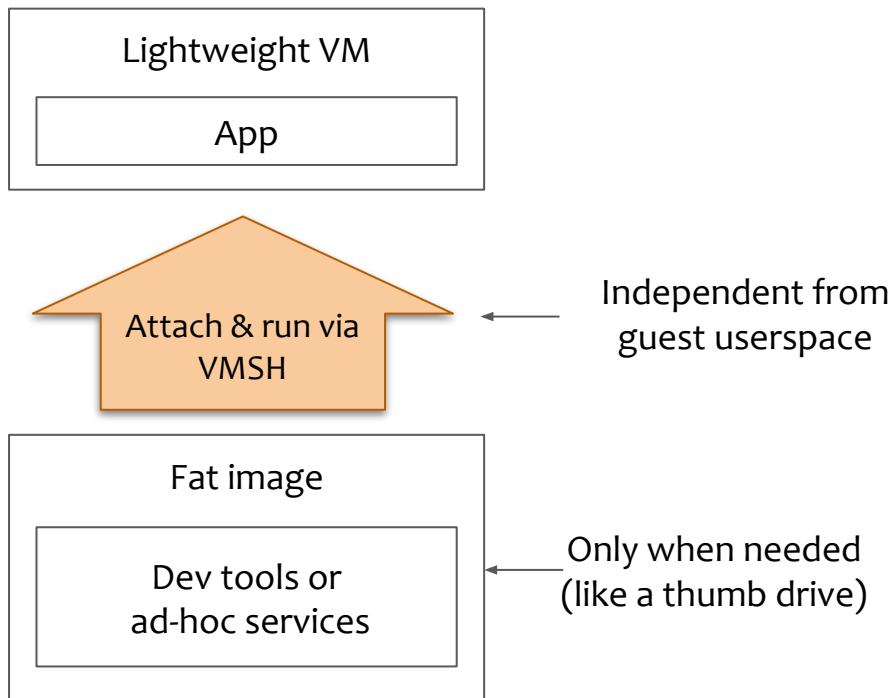


Out-of-band management with user-supplied tools?

VMSH: Guest overlays for VMs



VMSH: Guest overlays for VMs



VMSH attaches to VM **on demand** & without guest agents

Design


Design Goals Overview

Design goals

- Non-cooperativeness
 - No guest agents
- Generality
 - No hypervisor specific APIs
 - Many Linux kernels
- Performance
 - No degradation of guest processes

Overview

- Non-cooperativeness
 - Attach to any VM
- Generality
 - Side-load overlay container
- Performance
 - VMSH serves fat image



Implementation

Side-loading a kernel-agnostic library

Container-based system overlay


Side-loading a kernel-agnostic library

Side-loading:

- Side-load executable page into guest kernel
- Find kernel and parse its function table


The kernel library...

- Starts overlay container
- Starts VirtIO drivers



Container-based system overlay

- Overlay for attached tools
- Overlay with VMSH's block device as fs root
- Communication to outside world via VMSH devices
- VMSH VirtIO devices via ptrace and ioregionfd



Evaluation

Evaluation

Questions:

1. Is the implementation robust?
2. Is our approach general?
3. Does VMSH impact performance?

Experimental Testbed:

- Intel Core i9-9900K CPU
- 64GB RAM
- Intel P4600 NVMe 2TB

1. Is the implementation robust?

Xfstests [3]:

- File system testing
- Widely adopted by Linux devs
- Regression tests, fuzzing

Block device	Passing tests
Qemu	616
VMSH	616

VMSH's block device is as robust as Qemu's

2. Is our approach general?

4 KVM Hypervisors:


- ✓ QEMU
- ✓ kvmtool
- ✓ Firecracker
- ✓ crosVM

All Linux LTS kernels:

~40h to cover 5 years of kernel development


- ✓ v5.10, v5.4, v4.19, v4.14, v4.9, v4.4

3. Does VM SH impact performance?



3a. Common case: access original VM

3. Does VMSH impact performance?




3b. Attached tools: VMSH devices

3a. Overhead for the lightweight image

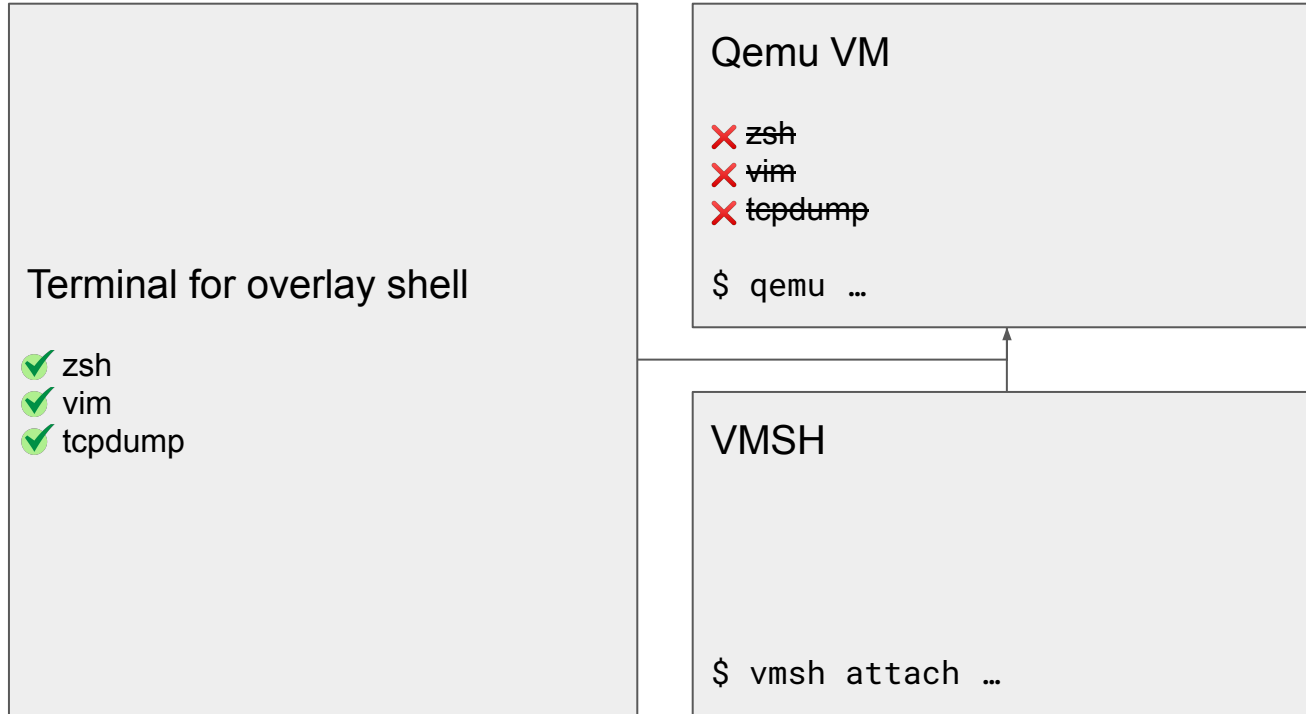
0x

For the common case of
accessing the original VM

3b. Overhead: VMESH devices



Demo



Conclusion

VMSH extends lightweight VMs with external functionality

- on-demand
- non-disruptively



VMSH provides...

1. A generic guest-overlay
2. Hypervisor-independent VirtIO devices
3. An OS-independent code side-loading into VM guests

Try it on <https://vmsh.org>

References



- [1] Maintenance icons created by kerismaker - Flaticon,
<https://www.flaticon.com/free-icons/maintenance>
- [2] Cube icons created by Freepik - Flaticon, <https://www.flaticon.com/free-icons/cube>
- [3] xfstests-dev <https://git.kernel.org/pub/scm/fs/xfs/xfstests-dev.git/>